

Actividad 4 - Pruebas de particionamiento de bases de datos NoSQL

Realizado por:

Joe Alejandro Sierra Ocassal

Ingrid Johana Rojas Gómez

Corporación Universitaria Iberoamericana

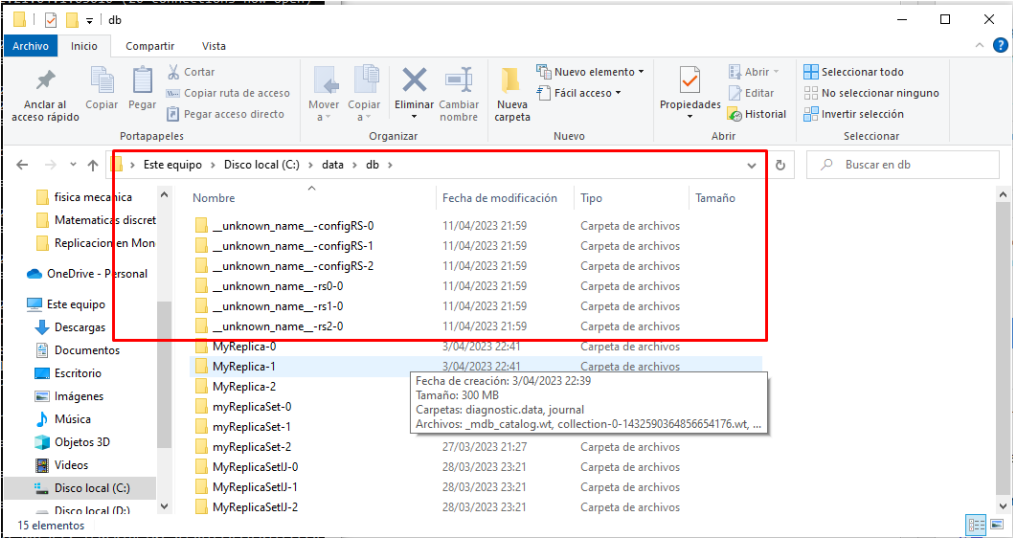
Bases de datos avanzadas

WILLIAM RUIZ

Abril 2023

Casos de prueba

Disponemos de los siguientes casos de prueba, en los cuales establecemos el objetivo de cada prueba a realizar, detallando el resultado obtenido, con el cual garantizamos la calidad y el funcionamiento del particionamiento dando cumplimiento a los requerimientos no funcionales documentados en la actividad #3.

#	Descripción	Resultado
1	<p>Particionamiento</p> <p>Se debe verificar que se crearon las particiones físicas correspondientes donde la información quede almacenada y a su vez se replica la data en cada shard como se especificó previamente en el esquema centralizado de replicas.</p>	<p>Se crea sin inconvenientes las particiones físicas:</p>  <p>Creación de los 3 shards:</p>

```

C:\Program Files\MongoDB\Server\4.2\bin\mongo.exe
{ "architecture": "x86_64", "version": "10.0 (build 19044)" } }
j20002| 2023-04-11T21:59:17.440-0500 I NETWORK [listener] connection accepted from 172.21.64.1:63087 #22 (11 connections now open)
j20002| 2023-04-11T21:59:17.441-0500 I NETWORK [conn22] received client metadata from 172.21.64.1:63087 conn22: { driver: { name: "NetworkInterfaceTL", version: "4.2.24" }, os: { type: "Windows", name: "Microsoft Windows 10", architecture: "x86_64", version: "10.0 (build 19044)" } }
j20002| 2023-04-11T21:59:17.569-0500 I CONNPPOOL [ReplicaSetMonitor-TaskExecutor] Connecting to DESKTOP-4TPBE5B:20001
j20002| 2023-04-11T21:59:17.569-0500 I CONNPPOOL [ReplicaSetMonitor-TaskExecutor] Connecting to DESKTOP-4TPBE5B:20000
j20002| 2023-04-11T21:59:17.569-0500 I NETWORK [shard-registry-reload] Starting new replica set monitor for __unknown_name__-rs2/DESKTOP-4TPBE5B:20002
j20002| 2023-04-11T21:59:17.571-0500 I CONNPPOOL [ReplicaSetMonitor-TaskExecutor] Connecting to DESKTOP-4TPBE5B:20002
j20002| 2023-04-11T21:59:17.574-0500 I NETWORK [ReplicaSetMonitor-TaskExecutor] Confirmed replica set for __unknown_name__-rs1 is __unknown_name__-rs1/DESKTOP-4TPBE5B:20001
j20002| 2023-04-11T21:59:17.574-0500 I SHARDING [Sharding-Fixed-1] Updating config server with confirmed set __unknown_name__-rs1/DESKTOP-4TPBE5B:20001
j20002| 2023-04-11T21:59:17.574-0500 I NETWORK [listener] connection accepted from 172.21.64.1:63090 #25 (12 connections now open)
j20002| 2023-04-11T21:59:17.574-0500 I NETWORK [ReplicaSetMonitor-TaskExecutor] Confirmed replica set for __unknown_name__-rs0 is __unknown_name__-rs0/DESKTOP-4TPBE5B:20000
j20002| 2023-04-11T21:59:17.574-0500 I SHARDING [Sharding-Fixed-1] Updating config server with confirmed set __unknown_name__-rs0/DESKTOP-4TPBE5B:20000
j20002| 2023-04-11T21:59:17.574-0500 I NETWORK [conn25] received client metadata from 172.21.64.1:63090 conn25: { driver: { name: "NetworkInterfaceTL", version: "4.2.24" }, os: { type: "Windows", name: "Microsoft Windows 10", architecture: "x86_64", version: "10.0 (build 19044)" } }
j20002| 2023-04-11T21:59:17.574-0500 I NETWORK [ReplicaSetMonitor-TaskExecutor] Confirmed replica set for __unknown_name__-rs2 is __unknown_name__-rs2/DESKTOP-4TPBE5B:20002
j20002| 2023-04-11T21:59:17.574-0500 I SHARDING [Sharding-Fixed-1] Updating config server with confirmed set __unknown_name__-rs2/DESKTOP-4TPBE5B:20002
s20006| 2023-04-11T21:59:16.185-0500 D1 SHARDING [shard-registry-reload] found 3 shards listed on config server(s) with lastVisibleOpTime: { ts: Timestamp(1681268349, 1), t: 1 }
s20006| 2023-04-11T21:59:16.185-0500 D1 NETWORK [shard-registry-reload] Started targeter for __unknown_name__-rs0/DESKTOP-4TPBE5B:20000
s20006| 2023-04-11T21:59:16.185-0500 D1 NETWORK [shard-registry-reload] Started targeter for __unknown_name__-rs1/DESKTOP-4TPBE5B:20001
s20006| 2023-04-11T21:59:16.185-0500 D1 NETWORK [shard-registry-reload] Started targeter for __unknown_name__-rs2/DESKTOP-4TPBE5B:20002
s20006| 2023-04-11T21:59:16.566-0500 D1 TRACKING [replSetDistLockPinger] Cmd: NotSet, TrackingId: 64361e8459ea17e2cc69a806
s20006| 2023-04-11T21:59:16.676-0500 D1 NETWORK [ReplicaSetMonitor-TaskExecutor] Refreshing replica set __unknown_name__-configRS took 0ms
s20006| 2023-04-11T21:59:16.791-0500 D1 TRACKING [UserCacheInvalidator] Cmd: NotSet, TrackingId: 64361e8459ea17e2cc69a808
s20006| 2023-04-11T21:59:17.270-0500 D1 EXECUTOR [ConfigServerCatalogCacheLoader-0] Reaping this thread; next thread reaped no earlier than 2023-04-11T21:59:47.270-0500
s20006| 2023-04-11T21:59:17.270-0500 D1 EXECUTOR [ConfigServerCatalogCacheLoader-0] shutting down thread in pool ConfigServerCatalogCacheLoader
s20006| 2023-04-11T21:59:17.419-0500 D1 NETWORK [ReplicaSetMonitor-TaskExecutor] Refreshing replica set __unknown_name__-rs0 took 0ms
s20006| 2023-04-11T21:59:17.538-0500 D1 NETWORK [ReplicaSetMonitor-TaskExecutor] Refreshing replica set __unknown_name__-rs1 took 0ms
s20006| 2023-04-11T21:59:17.654-0500 D1 NETWORK [ReplicaSetMonitor-TaskExecutor] Refreshing replica set __unknown_name__-rs2 took 0ms
s20006| 2023-04-11T21:59:17.659-0500 D1 EXECUTOR [UpdateReplicaSetOnConfigServer] Reaping this thread; next thread reaped no earlier than 2023-04-11T21:59:47.659-0500
s20006| 2023-04-11T21:59:17.659-0500 D1 EXECUTOR [UpdateReplicaSetOnConfigServer] shutting down thread in pool Sharding-Fixed
s20006| 2023-04-11T21:59:20.104-0500 D1 TRACKING [Uptime-reporter] Cmd: NotSet, TrackingId: 64361e8859ea17e2cc69a80a
    
```

Comprobación de registros insertados, donde se evidencia la replicación funcionando:

```

> primaryDB.deportistas.findOne()
{
  "_id" : ObjectId("6423b3b3fb308ae4113870ae"),
  "nombre" : "juan",
  "apellido" : "gonzales",
  "edad" : "26"
}

> secondaryDB.deportistas.findOne()
{
  "_id" : ObjectId("6423b3b3fb308ae4113870ae"),
  "nombre" : "juan",
  "apellido" : "gonzales",
  "edad" : "26"
}
    
```

		<pre>> thirdDB.deportistas.findOne() { "_id" : ObjectId("6423b3b3fb308ae4113870ae"), "nombre" : "juan", "apellido" : "gonzales", "edad" : "26" }</pre>
2	<p>Carga de datos</p> <p>Se debe validar el tiempo de inserción de datos, colección de Torneo Deportivo sea inferior a un segundo.</p>	<p>Para realizar la prueba a este caso, hemos intentado realizar la inserción de 800.000 datos a la colección de prueba Autores, en la que se insertaron 437.256 datos en 60 segundos:</p> <pre>> db = (new Mongo("localhost:20006")).getDB("Biblioteca") Biblioteca mongos> db.Autores.count() 437256 mongos></pre> <p>Por lo cual se encontró un rendimiento efectivo ya que por segundo se registraron 7.287 filas en la base de datos por segundo.</p>
3	<p>Tiempos de respuesta</p> <p>Validar tiempo de respuesta consulta de datos, alguna de las colecciones de la BD propuesto sea inferior a un segundo.</p>	<p>Realizamos la consulta de la colección específica de deportistas en la que consultamos con un límite de 20 registro con el fin de verificar su tiempo de respuesta óptimo con estos datos.</p> <pre>mongos> shard2DB.Deportistas.find().limit(20) { "_id" : ObjectId("6439f94458ed1561408448cd"), "id" : 0, "post_title" : "Blog Post by Author 0", "date" : ISODate("2023-04-15T01:09:24.872Z") } { "_id" : ObjectId("6439f94458ed1561408448ce"), "id" : 1, "post_title" : "Blog Post by Author 1", "date" : ISODate("2023-04-15T01:09:24.876Z") } { "_id" : ObjectId("6439f94458ed1561408448cf"), "id" : 2, "post_title" : "Blog Post by Author 2", "date" : ISODate("2023-04-15T01:09:24.877Z") } { "_id" : ObjectId("6439f94458ed1561408448d0"), "id" : 3, "post_title" : "Blog Post by Author 3", "date" : ISODate("2023-04-15T01:09:24.877Z") } { "_id" : ObjectId("6439f94458ed1561408448d1"), "id" : 4, "post_title" : "Blog Post by Author 4", "date" : ISODate("2023-04-15T01:09:24.878Z") } { "_id" : ObjectId("6439f94458ed1561408448d2"), "id" : 5, "post_title" : "Blog Post by Author 5", "date" : ISODate("2023-04-15T01:09:24.879Z") } { "_id" : ObjectId("6439f94458ed1561408448d3"), "id" : 6, "post_title" : "Blog Post by Author 6", "date" : ISODate("2023-04-15T01:09:24.879Z") } { "_id" : ObjectId("6439f94458ed1561408448d4"), "id" : 7, "post_title" : "Blog Post by Author 7", "date" : ISODate("2023-04-15T01:09:24.880Z") } { "_id" : ObjectId("6439f94458ed1561408448d5"), "id" : 8, "post_title" : "Blog Post by Author 8", "date" : ISODate("2023-04-15T01:09:24.881Z") } { "_id" : ObjectId("6439f94458ed1561408448d6"), "id" : 9, "post_title" : "Blog Post by Author 9", "date" : ISODate("2023-04-15T01:09:24.882Z") } { "_id" : ObjectId("6439f94458ed1561408448d7"), "id" : 10, "post_title" : "Blog Post by Author 10", "date" : ISODate("2023-04-15T01:09:24.882Z") } { "_id" : ObjectId("6439f94458ed1561408448d8"), "id" : 11, "post_title" : "Blog Post by Author 11", "date" : ISODate("2023-04-15T01:09:24.883Z") } { "_id" : ObjectId("6439f94458ed1561408448d9"), "id" : 12, "post_title" : "Blog Post by Author 12", "date" : ISODate("2023-04-15T01:09:24.884Z") } { "_id" : ObjectId("6439f94458ed1561408448da"), "id" : 13, "post_title" : "Blog Post by Author 13", "date" : ISODate("2023-04-15T01:09:24.884Z") } { "_id" : ObjectId("6439f94458ed1561408448db"), "id" : 14, "post_title" : "Blog Post by Author 14", "date" : ISODate("2023-04-15T01:09:24.885Z") } { "_id" : ObjectId("6439f94458ed1561408448dc"), "id" : 15, "post_title" : "Blog Post by Author 15", "date" : ISODate("2023-04-15T01:09:24.886Z") } { "_id" : ObjectId("6439f94458ed1561408448dd"), "id" : 16, "post_title" : "Blog Post by Author 16", "date" : ISODate("2023-04-15T01:09:24.887Z") } { "_id" : ObjectId("6439f94458ed1561408448de"), "id" : 17, "post_title" : "Blog Post by Author 17", "date" : ISODate("2023-04-15T01:09:24.887Z") } { "_id" : ObjectId("6439f94458ed1561408448df"), "id" : 18, "post_title" : "Blog Post by Author 18", "date" : ISODate("2023-04-15T01:09:24.888Z") } { "_id" : ObjectId("6439f94458ed1561408448e0"), "id" : 19, "post_title" : "Blog Post by Author 19", "date" : ISODate("2023-04-15T01:09:24.889Z") } mongos></pre> <p>Obtenemos un resultado exitoso en el cual tiene una respuesta inferior a 1 segundo.</p>
4	<p>Validación balanceo de carga de datos en los diferentes nodos:</p>	<p>Para este caso, hicimos la inserción de 800.000 datos con el balanceo de carga habilitado, de acuerdo con:</p>

Validar el funcionamiento del balanceador de carga de datos mediante los diferentes nodos creados, al hacer inserciones masivas de datos en alguna de las colecciones de la BD propuesto.

```

mongos> sh.status()
--- Sharding Status ---
  sharding version: {
    "_id" : 1,
    "minCompatibleVersion" : 5,
    "currentVersion" : 6,
    "clusterId" : ObjectId("6439f654d3d7061d4c3fc03b")
  }
  shards:
    { "_id" : "__unknown_name__-rs0", "host" : "__unknown_name__-rs0/DESKTOP-4TPBE5B:20000", "state" : 1 }
    { "_id" : "__unknown_name__-rs1", "host" : "__unknown_name__-rs1/DESKTOP-4TPBE5B:20001", "state" : 1 }
    { "_id" : "__unknown_name__-rs2", "host" : "__unknown_name__-rs2/DESKTOP-4TPBE5B:20002", "state" : 1 }
  active mongoses:
    "4.2.24" : 1
  autosplit:
    Currently enabled: yes
  balancer:
    Currently enabled: yes
    Currently running: no
    Failed balancer rounds in last 5 attempts: 0
    Migration Results for the last 24 hours:
      No recent migrations
  databases:
    { "_id" : "TorneoDeportivo", "primary" : "__unknown_name__-rs1", "partitioned" : true, "version" : { "uu
id" : UUID("ce831f1c-4fa8-4cda-a3e6-88820dcb44d7"), "lastMod" : 1 } }
      TorneoDeportivo.Deportistas
        shard key: { "id" : 1 }
        unique: false
        balancing: true
        chunks:
          __unknown_name__-rs1      1
          { "id" : { "$minKey" : 1 } } --> { "id" : { "$maxKey" : 1 } } on : __unknown_name__-rs1 Time
stamp(1, 0)
      { "_id" : "config", "primary" : "config", "partitioned" : true }
        config.system.sessions
          shard key: { "_id" : 1 }
          unique: false
          balancing: true
          chunks:
            __unknown_name__-rs0      1
            { "_id" : { "$minKey" : 1 } } --> { "_id" : { "$maxKey" : 1 } } on : __unknown_name__-rs0 Ti
mestamp(1, 0)

```

En la colección usada tenemos la siguiente verificación (después de la inserción de la cantidad de datos mencionados):

Shard2:

```

mongos> shard2DB.Deportistas.count()
400000

```

Shard1:

```

mongos> shard1DB.Deportistas.count()
400000

```

Por lo cual concluimos que se balancea la carga de datos.